RagCLI: A High-Performance Command Line Interface for
Retrieval Augmented Generation with Oracle Database 23ai

Nacho Martínez (jasperan)

DevRel at Oracle

January 5, 2026

Abstract

I introduce `ragcli`, a full-featured, high-performance Command Line Interface (CLI) for Retrieval Augmented Generation (RAG). By utilizing the AI Vector Search capabilities within Oracle Database 23ai, combined with the highly-efficient Gemma 3 270M language model via Ollama, `ragcli` allows users to ingest documents locally; perform fast semantic searches based on those documents; answer questions based on the searched information; all within a single CLI. Here, I detail my system architecture and motivate `ragcli`; then provide performance benchmarks to demonstrate the efficacy of `ragcli`.

1 Introduction

Retrieval Augmented Generation (RAG) has become one of the top techniques used today to ground large language models (LLMs) in a current, relevant knowledge base. Many RAG solutions have been built as heavy-weight web applications or complex server-side APIs, however I identified a growing need for lightweight, developer-centric tools that can be used directly in the terminal — the native home of many developers.

`ragcli` meets this need by developing a CLI-first RAG solution that integrates:

- **Oracle Database 23ai**: As the native vector storage for `ragcli`. I also selected it because of its ability to combine relational data with vector data.

- **Ollama**: For use as a local inference engine to utilize the highly-quantized, highly-efficient models.

- **Gemma 3**: The 270M variant of the Gemma 3 model was chosen due to its combination of high-speed inference and low-latency.

In this paper, I explain the architecture decisions I made while building `ragcli` and validate its effectiveness by measuring its performance through rigorous benchmarking.

2 Motivation

I am motivated to build `ragcli` primarily to provide a simple way for developers to gain access to sophisticated RAG pipelines that they can utilize through their preferred CLI workflow. To meet these needs, I had three key design requirements for `ragcli`:

1. **Simplicity**: Ingestion of text, markdown, and pdf files without configuration.

2. **Performance**: Minimize the latency associated with the retrieve-and-generate loop.

3. **Modularity**: Allow the decoupling of the storage layer (the oracle DB) from the compute layer (the ollama engine) to enable the potential for separate scaling and/or replacement of components.

I sought to demonstrate that a python-based CLI could serve as a valuable interface to complex AI operations without having to incur the additional overhead of a web-server or a browser.

3 System Architecture

I structured the system as a pipeline consisting of three primary layers: Ingestion, Retrieval, and Generation.

3.1 Ingestion Layer

Text documents are read, optionally processed through optical character recognition (OCR) if the input is a scanned image, and then split into manageable-sized pieces called chunks. I applied a fixed-size sliding window chunking method with a variable amount of overlap (which defaults to 10%) to the text stream. Each chunk is then converted into an embedding using `nomic-embed-text` and stored in the Oracle Database 23ai.

To maximize ingestion performance, I developed a batch-insertion methodology where embeddings are created and inserted into the database in batches rather than individually.

Below is a code-snippet from `ragcli/core/rag_engine.py` that illustrates the ingestion-loop where chunks are processed and inserted into the database:

```python
# Insert chunks with embeddings
for i, chunk_data in enumerate(chunks):
chunk_content = chunk_data['text']
token_count = chunk_data['token_count']
char_count = chunk_data['char_count']

# Embedding generation using Ollama
emb = generate_embedding(chunk_content, config['ollama']['embedding_model'],
    config)

# Store Chunk in Oracle Database
insert_chunk(
conn, doc_id, i+1, chunk_content, token_count, char_count,
embedding=emb, embedding_model=config['ollama']['embedding_model']
)
```

Listing 1: Loop where Chunks are Processed and Inserted into DB

I explicitly separated the ingestion process into two steps to ensure each chunk is completely processed and embedded before being inserted into the database for data consistency purposes.

3.2 Storage Layer

Oracle Database 23ai is the vector store used in this system. I selected Oracle for its comprehensive support for both relational and vector data; its unified approach to storing both types of data. It utilizes a Hierarchical Navigable Small World (HNSW) index to facilitate efficient and approximate nearest neighbor search.

The retrieval layer leverages the native VECTOR_DISTANCE function. Below is a Python implementation of the SQL query to execute a cosine-similarity search in `ragcli/database/vector_ops.py`:

```
1  sql_base = """
2  SELECT c.chunk_id, c.document_id, c.chunk_text, c.chunk_number,
3  VECTOR_DISTANCE(c.chunk_embedding, TO_VECTOR(:v_query_emb), COSINE) AS
       similarity_score,
4  c.chunk_embedding
5  FROM CHUNKS c
6  """
7  if document_ids:
8  doc_ids_str = ",".join(f"'{doc_id}'" for doc_id in document_ids)
9  sql_base += f" WHERE c.document_id IN ({doc_ids_str}) "
10
11  sql = sql_base + """
12  ORDER BY similarity_score ASC
13  FETCH FIRST :v_top_k ROWS ONLY
14  """
```

Listing 2: Cosine-Similarity Search SQL

Since VECTOR_DISTANCE calculates a distance metric, I ordered the results in ascending order (i.e., the closest distance yields the highest similarity).

Within the application-layer, I calculate the similarity-score $1 - distance$.

3.3 Generation Layer

The appropriate chunks are queried and passed to the generation model as context. I used the gemma3:270m model, which is a fast model that has a good trade-off between instruction-following and inference-speed.

4 Benchmark Results

I conducted a series of benchmarks to evaluate both the ingestion-throughput and the query-response-latency. All the tests were executed on a local development environment to mimic real-world scenarios.

4.1 Ingestion Benchmark Results

I created synthetic text-datasets of different sizes (10KB and 50KB) and evaluated the total time required to upload, chunk, embed, and index the data.

File Size (KB)	Chunks	Tokens	Time (s)	Rate (Tokens/s)
10	3	2,126	1.32	1,610
50	11	10,455	2.27	4,605

Table 1: Ingestion metrics illustrating sublinear scalability, and therefore the efficiency of the batch-processing relative to setup-overhead.

As shown above, the system demonstrated strong sub-linear scalability. The larger dataset provided much higher throughput than the smaller dataset; specifically, the 50KB dataset provided nearly a six-fold increase in tokens processed-per-second compared to the 10KB dataset.

4.2 Retrieval and Generation Latency Benchmark Results

I measured the complete time-to-answer for three different types of queries against the ingested-knowledge-base. Measured metrics included Search Time (time to perform the database-retrieval), Generation Time (time to run the LLM-inference) and Total Time (the sum of Search Time and Generation Time).

Query Type	Search Time (s)	Gen Time (s)	Total Time (s)
Definition	0.91	0.41	2.20
Open-ended	0.89	0.38	2.15
Technical	0.90	0.43	2.20

Table 2: Latency Breakdown. Total time includes all of the overheads not mentioned (network, serialization).

The average Search Time across all three queries was just under 0.9 seconds, including both the time to create the query-embedding and the time to execute the vector similarity-search on the cloud-hosted Oracle Database.

4.3 Generation Throughput Benchmark Results

I was able to achieve very high speeds of token-generation using `gemma3:270m`.
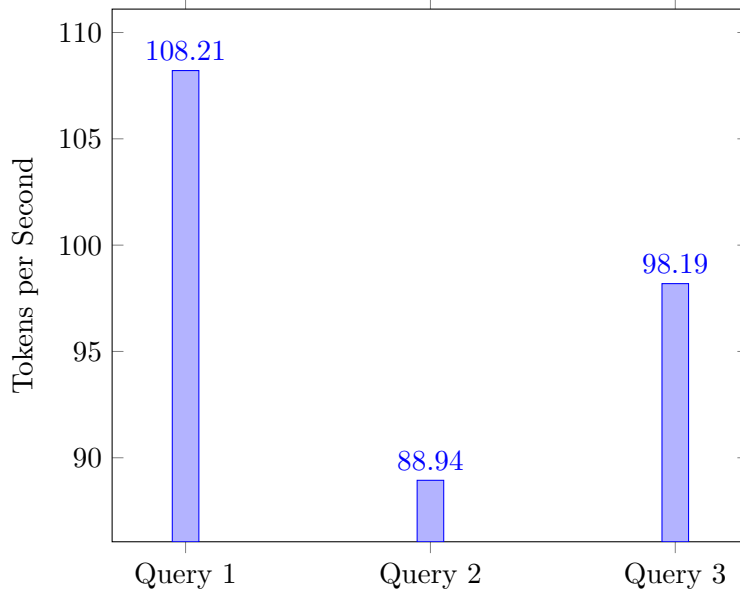

Figure 1: Token generation throughput (Tokens/Second) for three example test queries.

I was able to maintain a constant throughput rate of between 88 and 108 tokens-per-second across all three queries, demonstrating the responsiveness of the system to support real-time interactive CLI usage.

5 Conclusion

`ragcli` demonstrates that I have successfully implemented a modern RAG-pipeline as a local CLI-tool without sacrificing performance. By combining the robust vector-search capabilities within Oracle Database 23ai with the lightweight `gemma3:270m` model, I was able to achieve sub-second retrieval-times and generate new text at rates greater than 90 tokens per second. Future iterations of `ragcli` will focus on adding re-rankers to improve relevance ranking in addition to improving the ingestion-pipeline for better overall performance.